


Transformer Encoder-Decoder Architecture

Depicts the Transformer encoder-decoder stacks, residual sub-layers, and cross-attention connections from Figure 1.



Illustrative example. The Transformer connects sequence representations concurrently across N=6 stacked encoder and decoder layers, using masked self-attention and cross-attention. Simplified view omits sub-layer residual paths and internal head splits.