

The paper compares a frozen baseline with a small learned update.

An input reaches a frozen path and a trainable path, and their outputs are added. Adding the two outputs gives the layer output.

Two paths, one output

The frozen path and the learned update add to one output.

The two methods rank differently in the two evaluated settings.

For the toy input the frozen path returns 3 and the update returns 0.5, so the layer output is 3.5. The third drawing traces that computation step by step.

Step by step

The worked example returns 3.5.

Only two datasets were tested.